

Classification of wine: Supervised Machine Learning

Nigel K. Gondo

Brief Description

The purpose of this project is quite simple, it is to conduct a classification exercise to put wine into different classes of quality based on multiple variables such as acidity, sugar, alcohol etc.

Setting the directory with the data required

```
setwd('C://Users//Nigel Gondo//Documents//Portfolio  
Projects//Classification')
```

Importing the relevant libraries

```
#Importing Libraries  
library(dplyr)  
  
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##     filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##     intersect, setdiff, setequal, union  
  
library(caTools)  
library(randomForest)  
  
## randomForest 4.6-14  
  
## Type rfNews() to see new features/changes/bug fixes.  
  
##  
## Attaching package: 'randomForest'  
  
## The following object is masked from 'package:dplyr':  
##  
##     combine  
  
library(rpart)
```

Importing and Exploring the dataset

```
df_wine <- read.csv('WineQT.csv')
head(df_wine)
```

```
##   fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1           7.4           0.70         0.00           1.9       0.076
## 2           7.8           0.88         0.00           2.6       0.098
## 3           7.8           0.76         0.04           2.3       0.092
## 4          11.2           0.28         0.56           1.9       0.075
## 5           7.4           0.70         0.00           1.9       0.076
## 6           7.4           0.66         0.00           1.8       0.075
##   free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates alcohol
## 1                   11                   34 0.9978 3.51      0.56      9.4
## 2                   25                   67 0.9968 3.20      0.68      9.8
## 3                   15                   54 0.9970 3.26      0.65      9.8
## 4                   17                   60 0.9980 3.16      0.58      9.8
## 5                   11                   34 0.9978 3.51      0.56      9.4
## 6                   13                   40 0.9978 3.51      0.56      9.4
##   quality Id
## 1        5  0
## 2        5  1
## 3        5  2
## 4        6  3
## 5        5  4
## 6        5  5
```

#checking on the structure of the data frame

```
str(df_wine)
```

```
## 'data.frame':   1143 obs. of  13 variables:
## $ fixed.acidity      : num  7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 6.7 ...
## $ volatile.acidity   : num  0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58
0.58 ...
## $ citric.acid        : num  0 0 0.04 0.56 0 0 0.06 0 0.02 0.08 ...
## $ residual.sugar     : num  1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 1.8 ...
## $ chlorides          : num  0.076 0.098 0.092 0.075 0.076 0.075 0.069
0.065 0.073 0.097 ...
## $ free.sulfur.dioxide : num  11 25 15 17 11 13 15 15 9 15 ...
## $ total.sulfur.dioxide: num  34 67 54 60 34 40 59 21 18 65 ...
## $ density            : num  0.998 0.997 0.997 0.998 0.998 ...
## $ pH                 : num  3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36
3.28 ...
## $ sulphates          : num  0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57
0.54 ...
## $ alcohol            : num  9.4 9.8 9.8 9.8 9.4 9.4 9.4 10 9.5 9.2 ...
## $ quality            : int  5 5 5 6 5 5 5 7 7 5 ...
## $ Id                 : int  0 1 2 3 4 5 6 7 8 10 ...
```

#checking the dimensions of the data

```
dim(df_wine)
```

```
## [1] 1143    13

#quality needs to be converted into a categorical variable as this will the variable put into classes
df_wine$quality <- as.factor(df_wine$quality)
str(df_wine$quality)

## Factor w/ 6 levels "3","4","5","6",...: 3 3 3 4 3 3 3 5 5 3 ...

#checking if there are any null values in the data set
sum(is.na(df_wine))

## [1] 0

#Summary statistics
summary(df_wine)

## fixed.acidity    volatile.acidity    citric.acid    residual.sugar
## Min.   : 4.600    Min.   :0.1200    Min.   :0.0000    Min.   : 0.900
## 1st Qu.: 7.100    1st Qu.:0.3925    1st Qu.:0.0900    1st Qu.: 1.900
## Median : 7.900    Median :0.5200    Median :0.2500    Median : 2.200
## Mean   : 8.311    Mean   :0.5313    Mean   :0.2684    Mean   : 2.532
## 3rd Qu.: 9.100    3rd Qu.:0.6400    3rd Qu.:0.4200    3rd Qu.: 2.600
## Max.   :15.900    Max.   :1.5800    Max.   :1.0000    Max.   :15.500
## chlorides        free.sulfur.dioxide    total.sulfur.dioxide    density
## Min.   :0.01200    Min.   : 1.00        Min.   : 6.00        Min.   :0.9901
## 1st Qu.:0.07000    1st Qu.: 7.00        1st Qu.: 21.00       1st Qu.:0.9956
## Median :0.07900    Median :13.00       Median : 37.00       Median :0.9967
## Mean   :0.08693    Mean   :15.62       Mean   : 45.91       Mean   :0.9967
## 3rd Qu.:0.09000    3rd Qu.:21.00       3rd Qu.: 61.00       3rd Qu.:0.9978
## Max.   :0.61100    Max.   :68.00       Max.   :289.00       Max.   :1.0037
## pH              sulphates            alcohol            quality            Id
## Min.   :2.740    Min.   :0.3300    Min.   : 8.40    3: 6    Min.   : 0
## 1st Qu.:3.205    1st Qu.:0.5500    1st Qu.: 9.50    4: 33   1st Qu.: 411
## Median :3.310    Median :0.6200    Median :10.20    5:483   Median : 794
## Mean   :3.311    Mean   :0.6577    Mean   :10.44    6:462   Mean   : 805
## 3rd Qu.:3.400    3rd Qu.:0.7300    3rd Qu.:11.10    7:143   3rd Qu.:1210
## Max.   :4.010    Max.   :2.0000    Max.   :14.90    8: 16   Max.   :1597
```

Splitting the dataset into train and test to train and predict the model


```
set.seed(42)
split_df_wine <- sample.split(df_wine, SplitRatio = 0.8)
training_set <- subset(df_wine, split_df_wine == TRUE)
testing_set <- subset(df_wine, split_df_wine == FALSE)

#Creating the random forest model and displaying the metrics
rf <- randomForest(quality~.,data=training_set)
rf
```


```
##
## Call:
## randomForest(formula = quality ~ ., data = training_set)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 3
##
##           OOB estimate of  error rate: 30.11%
## Confusion matrix:
##   3 4  5  6  7 8 class.error
## 3 0 0  4  0 0 0  1.0000000
## 4 0 0 20  5 0 0  1.0000000
## 5 0 0 303 78 2 0  0.2088773
## 6 0 0 82 255 15 1  0.2776204
## 7 0 0  2 41 57 2  0.4411765
## 8 0 0  0  7 6 0  1.0000000
```

The error rate of the model is 30.11%, meaning that the accuracy of the model is 69.89%

```
#Variable importance
varImpPlot(rf)
```



```
plot(rf)
```



The error rate is stabilising as more trees are added